

25. Semaphores



semaphore : a single primitive for all things related to synchronization
one can use semaphore as both locks and condition variables.

Semaphores : a Definition

semaphore : object with an integer value that we can manipulate with two routines

- `sem_wait()` / `sem_post()`

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1);
```

Figure 31.1: Initializing A Semaphore

- 2nd : 0 , semaphore is shared between threads in the same process
- 3rd : 1 , 1 semaphore

```
1  int sem_wait(sem_t *s) {
2      decrement the value of semaphore s by one
3      wait if value of semaphore s is negative
4  }
5
6  int sem_post(sem_t *s) {
7      increment the value of semaphore s by one
8      if there are one or more threads waiting, wake one
9  }
```

Figure 31.2: Semaphore: Definitions Of Wait And Post

Binary Semaphores (Locks)

```
1  sem_t m;
2  sem_init(&m, 0, X); // init to X; what should X be?:1
3
4  sem_wait(&m);
5  // critical section here
6  sem_post(&m);
```

Figure 31.3: A Binary Semaphore (That Is, A Lock)

Setting the Value of a Semaphore

consider the number of resources you are willing to give away immediately after initialization

Semaphores For Ordering

```
1  sem_t s;
2
3  void *child(void *arg) {
4      printf("child\n");
5      sem_post(&s); // signal here: child is done
6      return NULL;
7  }
8
9  int main(int argc, char *argv[]) {
10     sem_init(&s, 0, X); // what should X be?    :0
11     printf("parent: begin\n");
12     pthread_t c;
13     Pthread_create(&c, NULL, child, NULL);
14     sem_wait(&s); // wait here for child
15     printf("parent: end\n");
16     return 0;
17 }
```

Figure 31.6: A Parent Waiting For Its Child

using a semaphore as an ordering primitive (similar to use of condition variables)

The Producer / Consumer (Bounded Buffer) Problem

First Attempt

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4
5  void put(int value) {
6      buffer[fill] = value;    // Line F1
7      fill = (fill + 1) % MAX; // Line F2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // Line G1
12     use = (use + 1) % MAX;    // Line G2
13     return tmp;
14 }
```

Figure 31.9: The Put And Get Routines

what if 2 threads run put() at the same time before one increase fill? : overwritten

```

1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);          // Line P1
8          put(i);                    // Line P2
9          sem_post(&full);            // Line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);             // Line C1
17         tmp = get();                 // Line C2
18         sem_post(&empty);            // Line C3
19         printf("%d\n", tmp);
20     }
21 }
22
23 int main(int argc, char *argv[]) {
24     // ...
25     sem_init(&empty, 0, MAX); // MAX are empty
26     sem_init(&full, 0, 0);    // 0 are full
27     // ...
28 }

```

Figure 31.10: Adding The Full And Empty Conditions

A solution: adding mutual exclusion

```

1  void *producer(void *arg) {
2      int i;
3      for (i = 0; i < loops; i++) {
4          sem_wait(&mutex);           // Line P0 (NEW LINE)
5          sem_wait(&empty);           // Line P1
6          put(i);                     // Line P2
7          sem_post(&full);             // Line P3
8          sem_post(&mutex);           // Line P4 (NEW LINE)
9      }
10 }
11
12 void *consumer(void *arg) {
13     int i;
14     for (i = 0; i < loops; i++) {
15         sem_wait(&mutex);             // Line C0 (NEW LINE)
16         sem_wait(&full);              // Line C1
17         int tmp = get();              // Line C2
18         sem_post(&empty);             // Line C3
19         sem_post(&mutex);             // Line C4 (NEW LINE)
20         printf("%d\n", tmp);
21     }
22 }

```

Deadlock
: consumer holds lock first
=> wait forever

Figure 31.11: Adding Mutual Exclusion (Incorrectly)

At Last, a Working Solution

```
1 void *producer(void *arg) {
2     int i;
3     for (i = 0; i < loops; i++) {
4         sem_wait(&empty);           // Line P1
5         sem_wait(&mutex);           // Line P1.5 (lock)
6         put(i);                     // Line P2
7         sem_post(&mutex);           // Line P2.5 (unlock)
8         sem_post(&full);            // Line P3
9     }
10 }
11
12 void *consumer(void *arg) {
13     int i;
14     for (i = 0; i < loops; i++) {
15         sem_wait(&full);            // Line C1
16         sem_wait(&mutex);           // Line C1.5 (lock)
17         int tmp = get();            // Line C2
18         sem_post(&mutex);           // Line C2.5 (unlock)
19         sem_post(&empty);           // Line C3
20         printf("%d\n", tmp);
21     }
22 }
```

mutex acquire and release to be just around the critical section
*better put that in the put() & get() function

Figure 31.12: Adding Mutual Exclusion (Correctly)

Simple and Dum can be better

with locking, sometimes a simple spin lock works best, because it is easy to implement and fast. Complex can mean slow. => Always try the simple & dumb approach first

Reader-Writer Locks

different data structure accesses might require different kinds of locking
as long as we can guarantee that no insert is on-going, we can allow many lookups to proceed concurrently.

```
1  typedef struct _rwlock_t {
2      sem_t lock;          // binary semaphore (basic lock)
3      sem_t writelock;     // allow ONE writer/MANY readers
4      int    readers;      // #readers in critical section
5  } rwlock_t;
6
7  void rwlock_init(rwlock_t *rw) {
8      rw->readers = 0;
9      sem_init(&rw->lock, 0, 1);
10     sem_init(&rw->writelock, 0, 1);
11 }
12
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     rw->readers++;
16     if (rw->readers == 1) // first reader gets writelock
17         sem_wait(&rw->writelock);
18     sem_post(&rw->lock);
19 }
20
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0) // last reader lets it go
25         sem_post(&rw->writelock);
26     sem_post(&rw->lock);
27 }
28
29 void rwlock_acquire_writelock(rwlock_t *rw) {
30     sem_wait(&rw->writelock);
31 }
32
33 void rwlock_release_writelock(rwlock_t *rw) {
34     sem_post(&rw->writelock);
35 }
```


Figure 31.13: A Simple Reader-Writer Lock

- Once a reader has acquired a read lock, more readers will be allowed to acquire the read lock too.
- any thread that wishes to acquire the write lock will have to wait until all readers are finished
- the last one to exit the critical section calls `sem_post()` on “writelock” and thus enables a waiting writer to acquire the lock.

Negatives : fairness. It would be relatively easy for readers to starve writers.

Reader-writer locks often add more overhead (especially with more sophisticated implementations)
=> thus do not end up speeding up performance as compared to just using simple and fast locking primitives

The Dining Philosophers



```
int left(int p) { return p; }  
int right(int p) { return (p + 1) % 5; }
```

Figure 31.14: The Dining Philosophers

Broken solution

```
1 void get_forks(int p) {  
2     sem_wait(&forks[left(p)]);  
3     sem_wait(&forks[right(p)]);  
4 }  
5  
6 void put_forks(int p) {  
7     sem_post(&forks[left(p)]);  
8     sem_post(&forks[right(p)]);  
9 }
```

what if everyone grab their
left fork? => deadlock

Figure 31.15: The `get_forks()` And `put_forks()` Routines

A solution : Breaking the Dependency

change how forks are acquired by at least one of the philosophers

: last philosopher grabs right fork before left.

```
1 void get_forks(int p) {  
2     if (p == 4) {  
3         sem_wait(&forks[right(p)]);  
4         sem_wait(&forks[left(p)]);  
5     } else {  
6         sem_wait(&forks[left(p)]);  
7         sem_wait(&forks[right(p)]);  
8     }  
9 }
```

Figure 31.16: Breaking The Dependency In `get_forks()`

Thread Throttling (admission control)

how can a programmer prevent “too many” threads from doing something at once and bogging the system down? (ex. All threads enter the memory-intensive region at the same time)

: decide a threshold for “too many”

=> use a semaphore to limit the number of threads concurrently executing.

How to implement Semaphores

```
1  typedef struct __Zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  } Zem_t;
6
7  // only one thread can call this
8  void Zem_init(Zem_t *s, int value) {
9      s->value = value;
10     Cond_init(&s->cond);
11     Mutex_init(&s->lock);
12 }
13
14 void Zem_wait(Zem_t *s) {
15     Mutex_lock(&s->lock);
16     while (s->value <= 0)
17         Cond_wait(&s->cond, &s->lock);
18     s->value--;
19     Mutex_unlock(&s->lock);
20 }
21
22 void Zem_post(Zem_t *s) {
23     Mutex_lock(&s->lock);
24     s->value++;
25     Cond_signal(&s->cond);
26     Mutex_unlock(&s->lock);
27 }
```

Figure 31.17: Implementing Semaphores With Locks And CVs

building condition variables out of semaphores is a much trickier proposition

Semaphore wait/post can only “hint” that something happened, but they can’t directly enforce *when* or *under what condition* a thread should wake without extra machinery.

Be Careful with Generalization

: Don’t generalize. Generalizations are generally wrong.

One could view semaphores as a generalization of locks and condition variables.

However, given the difficulty of realizing a condition variable on top of a semaphore, perhaps this generalization is not as general as you might think.